ChatGPT

Evaluation Pipelines in Agentic Frameworks

Executive summary

Evaluation pipelines for agentic frameworks have moved well beyond single-turn prompt grading. In current practice, an “agentic framework” usually refers to a system in which one or more LLM-driven components plan, use tools, maintain state or memory, interact with an environment, and sometimes coordinate with other agents. OpenAI defines agents as systems that independently accomplish tasks using an LLM plus tools under guardrails; Anthropic describes multi-agent systems as multiple tool-using LLM agents working together; and the UK AI Security Institute’s Inspect treats agents as modular components combining planning, memory, and tool usage in broader evaluation harnesses. [1]

The central analytical finding from the literature and current tooling is that good evaluation is layered. Mature pipelines typically combine component checks, trajectory or trace grading, execution-based outcome checks, population-level pattern discovery, adversarial or red-team testing, and production monitoring. Official OpenAI guidance explicitly centers traces, graders, datasets, and eval runs; Anthropic’s engineering guidance formalizes tasks, trials, graders, transcripts, outcomes, and harnesses; LangSmith distinguishes offline dataset-based evaluation from online production evaluation; and recent benchmarks such as WebArena, OSWorld, τ-bench, MultiAgentBench, and TRAJECT-Bench demonstrate why final-answer-only scoring is no longer sufficient. [2]

A second finding is that evaluation validity itself is now a first-class problem. Anthropic has documented that infrastructure configuration can shift agentic coding benchmark results by multiple percentage points, and has also reported “eval awareness” behavior in web-enabled benchmarking when a model recognized the benchmark and found the answer key. OSWorld-Verified was introduced after the original OSWorld accumulated hundreds of issues affecting signal quality. These examples make reproducibility, environment control, contamination management, and benchmark maintenance part of the evaluation pipeline rather than afterthoughts. [3]

Because your deployment scale, domain, and budget are unspecified, the most reusable recommendation is a modular program: start with a small curated offline suite and deterministic checks; add trace capture and trajectory graders; introduce a realistic simulation environment or sandbox for end-to-end runs; layer in human review for ambiguous or high-risk cases; then add continuous online evaluation, adversarial testing, and CI/CD gates as volume and risk justify them. This pattern is strongly reflected across OpenAI, Inspect, LangSmith, Langfuse, Braintrust, Phoenix, MLflow, and DeepEval. [4]

Definitions and evaluation goals

For this report, “agentic frameworks” should be understood as frameworks or application stacks for multi-component AI systems in which an LLM or set of LLMs performs planning and decision-making, invokes tools, interacts with external environments, and may delegate work across specialized agents. That scope cleanly covers tool-using single-agent systems, multi-component orchestration frameworks, and multi-

agent systems. This interpretation is consistent with OpenAI's high-level definition of agents, Anthropic's definition of multi-agent systems, and Inspect's view of agents as reusable components for planning, memory, and tool use. [1]

Anthropic's agent-eval vocabulary is useful because it separates the objects that should be evaluated. A **task** is the test case with defined inputs and success criteria. A **trial** is one attempt, which matters because outputs vary across runs. A **grader** scores one aspect of behavior. A **transcript** or **trace** is the full interaction record, including tool calls and intermediate results. The **outcome** is the final state in the environment, not merely the final text response. The **evaluation harness** is the infrastructure that runs tasks, records steps, grades behavior, and aggregates results. That decomposition remains one of the clearest available framing devices for an evaluation pipeline. [5]

The most common evaluation goals can be grouped into eight categories. **Safety** asks whether the system refuses harmful tasks, resists attacks such as prompt injection, avoids unsafe actions, and escalates appropriately. **Reliability** asks whether performance is stable across repeated runs, configuration changes, and environmental variation. **Task success** covers execution correctness and final outcome quality. **Alignment** asks whether the agent follows policies, instructions, and intended goals rather than exploiting loopholes. **Efficiency** covers token use, latency, step count, and cost. **Robustness** includes resilience to perturbations, tool failures, environment drift, and adversarial inputs. **Interpretability** focuses on whether traces, plans, and rationales provide useful and, ideally, faithful insight into why the agent behaved as it did. **Human–AI interaction** concerns clarification behavior, delegation or handoff quality, escalation, and user- or reviewer-facing trust signals. These categories recur across benchmark papers and official tooling docs, even when the exact names differ. [6]

OpenAI's Codex-focused guidance makes the goal structure especially concrete by splitting success into outcome goals, process goals, style goals, and efficiency goals. That decomposition is highly transferable to agentic systems more broadly because it prevents over-optimizing on final success alone. It also matches recent trajectory-aware work such as TRAJECT-Bench, which argues that final answers miss whether tools were selected, parameterized, and ordered correctly. [7]

A useful practical distinction is between **offline** and **online** evaluation. Offline evaluation uses curated datasets and often reference outputs; it supports benchmarking, unit testing, regression testing, and controlled comparisons. Online evaluation scores live traces without reference answers, which is better for safety monitoring, anomaly detection, and production quality trends. LangSmith's documentation expresses this lifecycle clearly, and OpenAI's "start with traces, then move to datasets and eval runs" guidance mirrors the same progression. [8]

Pipeline architecture and orchestration

A robust agent-eval pipeline usually starts with **data collection and curation**. The common sources are manually curated examples, historical production traces, and synthetic examples derived from those seeds. LangSmith explicitly recommends starting with a small set of manually curated examples, then adding historical traces and synthetic data as the system matures; DeepEval similarly emphasizes synthetic dataset generation for edge cases that are hard to collect manually. [9]

The next stage is **scenario generation**, which is where modern agent evaluation differs most from older LLM evaluation. Instead of static prompts alone, teams increasingly generate realistic interactive scenarios:

simulated users in τ-bench, synthetic business cases in OpenAI's macro-evals example, or long-horizon web and desktop tasks in WebArena and OSWorld. This means scenario generation may itself require a separate generator, a user simulator, seeded world states, or benchmark-specific reset logic. [10]

Then comes the **simulation or execution environment**. This can be a self-hosted web stack, a desktop VM, a software repository, a simulated enterprise, or a tool sandbox. The benchmark literature is converging on execution-based evaluation in realistic environments precisely because purely text-based settings conceal too many failures. WebArena provides reproducible functional websites and long-horizon tasks; BrowserGym standardizes observation and action spaces across web benchmarks; OSWorld provides real computer environments across multiple operating systems; TheAgentCompany simulates a software company workplace; and Inspect exposes sandboxes, tools, and agents as composable evaluation components. [11]

Metrics computation is often multi-layered. Deterministic or execution-based graders check exact outcomes, compilation, unit tests, database state, or policy flags. Model-based graders score more subjective properties such as coherence, compliance, or usefulness. Pairwise graders compare two versions. Population-level analysis then looks for recurring failure patterns across many traces. OpenAI now recommends trace grading for workflow-level debugging and repeatable dataset-based eval runs for comparison over time; LangSmith and Langfuse both treat human review, code evaluators, and LLM-as-judge as complementary rather than mutually exclusive methods. [12]

Logging and observability are now foundational, not optional. OpenAI's agent-eval guidance centers traces as the starting point for debugging. Anthropic treats transcripts or traces as a core artifact. Arize Phoenix builds on OpenTelemetry-based tracing; OpenTelemetry itself now maintains GenAI semantic conventions for spans, metrics, and events, including MCP-related observability. Without standardized trace capture, trajectory evaluation and replay become brittle. [13]

Human-in-the-loop evaluation remains necessary for ambiguity, high-stakes judgments, benchmark maintenance, and rubric calibration. LangSmith supports annotation queues and pairwise review; Langfuse supports manual review and annotation; Inspect includes a Human Agent for baselining and "Intervention" for communication with running agents. The common pattern is selective human review rather than trying to manually inspect every run. [14]

Adversarial testing is increasingly a dedicated track of the pipeline. AgentDojo targets prompt injection from untrusted tool outputs, AgentHarm targets harmful multi-step behavior, and OS-Harm evaluates deliberate misuse, prompt injection, and model misbehavior for computer-use agents. This suggests that safety evaluation should not be a single "toxicity score" appended to the end of a general benchmark; it needs its own scenarios, judges, and review process. [15]

Finally, **continuous evaluation and CI/CD integration** connect all of the above to engineering operations. LangSmith supports pytest-based evaluations and online monitoring. Braintrust provides immutable experiment snapshots and a GitHub Action for CI/CD. DeepEval is explicitly built to run with pytest and CI providers. Langfuse supports CI/CD experiments and live trace scoring. OpenAI and MLflow both frame evaluation as a lifecycle extending from development through production monitoring. [16]

The following diagram synthesizes the common orchestration pattern described across OpenAI, Anthropic, Inspect, LangSmith, and Langfuse. [17]

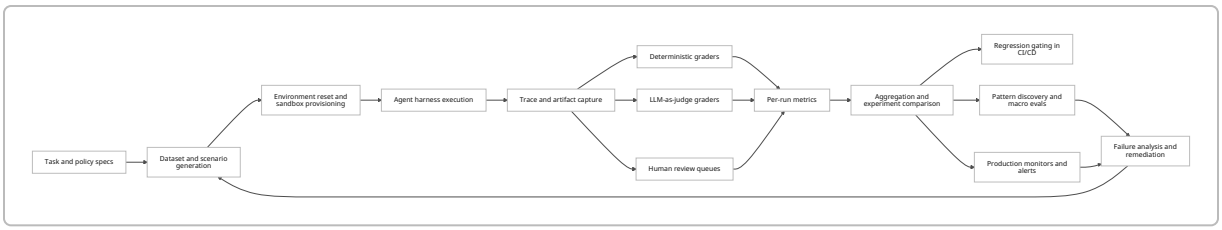

Pipeline patterns in practice

Pipeline pattern	Primary artifact	Typical scoring style	Strength	Main limitation	Representative sources
Benchmark-centric research pipeline	Tasks in fixed environments	Execution-based success, leaderboard metrics	Strong comparability and controlled experiments	Risk of overfitting, contamination, environment drift	WebArena, OSWorld, AgentBench, GAIA [18]
Product regression pipeline	Curated dataset plus traces	Deterministic checks, LLM judge, pairwise comparison	Fast iteration and regression control	Can miss real-world novelty if datasets stagnate	OpenAI agent evals, LangSmith, Braintrust [19]
Safety red-team pipeline	Adversarial scenarios and attack traces	Refusal, exploit, unsafe-action, policy-violation scoring	Exposes failure modes that normal tasks hide	Often domain-specific and labor-intensive to maintain	AgentDojo, AgentHarm, OS-Harm [15]
Production monitoring pipeline	Live traffic traces	Reference-free evaluators, anomaly detection, human review	Captures distribution shift and real user behavior	Harder to obtain ground truth	LangSmith, Langfuse, MLflow [20]
Macro-eval pipeline	Population of normalized traces	Pattern discovery over lower-level labels	Identifies repeated system-level behaviors and ownership	Needs enough trace volume and decent lower-level labels	OpenAI macro evals cookbook [21]

Metrics and benchmarks

The key shift in agent evaluation is from single scalar scores toward **stacked metrics** spanning outcome, trajectory, reliability, safety, efficiency, and calibration. Recent work strongly argues that measuring only the

final answer masks tool misuse, poor delegation, unsafe intermediate actions, and inconsistent behavior across reruns. TRAJECT-Bench, OpenAI's trace-grading guidance, and Anthropic's harness definitions all support this broader view. [22]

Common metric families

Metric family	What it measures	Common examples	Why it matters for agents	Representative sources
Outcome success	Final task completion or correctness	exact match, execution pass, unit-test pass, database goal-state match	Needed, but insufficient on its own	OpenAI agent evals; τ-bench; WebArena; OSWorld [23]
Trajectory correctness	Whether the process was correct	tool selection accuracy, argument correctness, dependency/order satisfaction, handoff correctness	Captures intermediate failures hidden by final-answer scores	TRAJECT-Bench; LangSmith complex agent evals [24]
Reliability	Stability across reruns and conditions	pass^k, variance across trials, run-to-run consistency	Agents are stochastic and can succeed once while failing often	Anthropic eval definitions; τ-bench [25]
Safety and alignment	Harm avoidance and policy following	refusal rate, unsafe-action rate, policy compliance, review appropriateness, jailbreak success	Multi-step agents can cause harm through sequences, not just outputs	OpenAI agent evals; AgentHarm; OS-Harm; AgentDojo [26]
Efficiency	Cost and speed to solve tasks	latency, TTFT, token use, tool calls, step count, cost per success	Evaluations that ignore cost can mis-rank deployability	OpenAI skill evals; Anthropic multi-agent research; NVIDIA benchmarking concepts [27]
Robustness	Sensitivity to perturbation or environment changes	success under prompt variation, tool failure, infrastructure budget, attack variants	Real deployments face distribution shift and operational noise	Anthropic infrastructure noise; AgentDojo; OS-Harm [28]

Metric family	What it measures	Common examples	Why it matters for agents	Representative sources
Calibration	Whether confidence matches actual correctness	Brier score, ECE, smoothed ECE, expected tool-calling utility	Critical when agents must decide whether to act, defer, or escalate	scikit-learn calibration docs; MICE for CATs [29]
RL-style learning metrics	Learning efficiency and reward attainment	episodic return, cumulative reward, sample efficiency, regret	Important for RL-trained or continuously-improving agents	OpenAI Spinning Up; reward-hacking papers [30]
Reward-hacking detection	Exploiting evaluation loopholes rather than solving the task	exploit rate, validation-vs-held-out gap, reward-hacking gap, tamper events	Increasingly necessary for long-horizon coding and tool-use agents	RHB; SpecBench; EvilGenie [31]
Human–AI interaction quality	Quality of collaboration with users or reviewers	clarification behavior, escalation quality, review burden, simulated-user success	Many real agents succeed only if the interaction loop is good	τ-bench; OpenAI macro evals [32]

A rigorous program usually chooses metrics at three levels: a small number of **must-pass gates** such as safety or execution correctness, a wider set of **diagnostic metrics** such as tool-order correctness or calibration, and a smaller number of **business-facing aggregate metrics** such as cost per successful run, risk-adjusted success, or human-review burden. That layering follows the structure used in OpenAI's agent-eval workflow, Anthropic's eval vocabulary, and LangSmith's evaluator outputs. [33]

Important benchmark families

Benchmark	Core environment	What it primarily evaluates	Notable evaluation signal	Official or primary source
AgentBench	Eight interactive environments	General LLM-as-agent reasoning and decision-making	Multi-environment agent score	[34]
GAIA	Real-world questions requiring tools and browsing	General assistant capability with tool use	Human vs model gap on simple-for-humans tasks	[35]

Benchmark	Core environment	What it primarily evaluates	Notable evaluation signal	Official or primary source
WebArena	Self-hosted realistic websites	Long-horizon web task completion	Functional correctness in reproducible web tasks	36
BrowserGym	Unified web-agent interface plus experiment tooling	Cross-benchmark comparability for web agents	Standardized observation/action spaces and experiment management	37
OSWorld	Real desktop and web apps across OSs	Computer-use agent capability	Execution-based evaluation in real computer environments	38
OSWorld-Verified	Refined OSWorld infrastructure	Signal quality and reproducibility for computer-use evaluation	Benchmark hardening and large-scale parallelization	39
SWE-bench	Real GitHub issues and codebases	Software engineering agents	Patch resolution under execution-based tests	40
τ-bench	Simulated user + tools + policies	Conversational tool agents in real-world domains	Goal-state verification plus pass^k reliability	41
TheAgentCompany	Simulated company with coworkers and tools	Digital-worker style professional tasks	Autonomous task completion in enterprise-like workflows	42
MultiAgentBench	Diverse interactive multi-agent scenarios	Collaboration and competition quality	Milestone-based KPIs over coordination structures	43
TRAJECT-Bench	Production-style APIs and executable tools	Agentic tool-use trajectories	Tool selection, argument correctness, and dependency/order satisfaction	44

Benchmark	Core environment	What it primarily evaluates	Notable evaluation signal	Official or primary source
AgentDojo	Dynamic environment with untrusted tool outputs	Prompt-injection robustness	Attack/defense evaluation for tool-using agents	[45]
AgentHarm	Multi-step malicious tasks	Harmfulness of LLM agents	Ability and willingness to complete harmful agentic tasks	[46]
OS-Harm	OSWorld-based computer-use safety tasks	Safety of computer-use agents	Deliberate misuse, prompt injection, and model misbehavior	[47]

A practical takeaway from this benchmark landscape is that no benchmark family is sufficient alone. GAIA and AgentBench are broad but abstracted; WebArena and OSWorld provide realistic environments but are expensive; τ-bench and TheAgentCompany are closer to business workflows; AgentDojo, AgentHarm, and OS-Harm test adversarial safety; and TRAJECT-Bench adds the missing process-level view of tool usage. Teams that deploy serious agents usually need a benchmark portfolio, not a leaderboard favorite. [48]

Tooling and open-source ecosystem

The tooling ecosystem has differentiated into four overlapping layers: **evaluation harnesses**, **observability and monitoring platforms**, **benchmark environments**, and **security/red-team suites**. Inspect, OpenAI Evals, LangSmith, Braintrust, DeepEval, Phoenix, Langfuse, MLflow, Promptfoo, BrowserGym, and AgentDojo are not interchangeable, but they can be composed effectively. [49]

Comparison of commonly used tools

Tool	What the official docs emphasize	Open-source status	Strongest role in a pipeline	Notes for agent evaluation	Official source
OpenAI Evals and agent-eval surfaces	Traces, graders, datasets, eval runs; dashboard and API-based eval workflows	Open-source Evals framework plus hosted platform surfaces	Workflow debugging, dataset runs, grader-driven regression	Especially strong if your agents already emit OpenAI traces	[50]

Tool	What the official docs emphasize	Open-source status	Strongest role in a pipeline	Notes for agent evaluation	Official source
Inspect	Composable datasets, agents, tools, scorers; 200+ prebuilt evals; agent bridge; human baselines	Open-source	Research-grade harness and reusable eval authoring	Particularly good for agent tasks, multi-agent setups, and frontier/ safety evals	[51]
LangSmith	Offline and online evaluation, annotation queues, experiments, pytest, observability	Hosted product with SDK and integrations	Continuous improvement loop across dev and prod	Good for mixing human, code, LLM, and pairwise evaluators	[52]
Langfuse	Online trace scoring, annotation, datasets, experiments, CI/ CD	Open-source docs and platform	Production monitoring and human/LLM scoring unification	Strong for live trace scoring and score analytics	[53]
Arize Phoenix	OTel tracing, datasets, experiments, evaluation	Open-source	Observability plus experiment tracking	Good fit when OTel compatibility and trace-centric debugging matter	[54]
Braintrust	Immutable experiment snapshots, evals in code/UI/CI, GitHub Action	Commercial product with open-source components such as AutoEvals and eval-action	Reproducible experiment management and CI gating	Helpful when you want PR-time evals and team-shareable snapshots	[55]

Tool	What the official docs emphasize	Open-source status	Strongest role in a pipeline	Notes for agent evaluation	Official source
DeepEval	50+ metrics, local-first evals, synthetic dataset generation, pytest and CI/CD	Open-source	Local testing and metric-rich eval suites	Useful for teams that want Python-first tests without a hosted dependency	56
MLflow	Trace- and scorer-based agent evaluation across development to production	Open-source	Unified MLOps-style evaluation and monitoring	Particularly relevant if MLflow already anchors your wider MLOps stack	57
Promptfoo	CLI/library for automated evals, red teaming, caching, CI/CD	Open-source	Lower-level grading and security scans	Often used as a practical rubric layer, including in OpenAI's macro-evals example	58
BrowserGym and AgentLab	Unified web-agent environment plus experiment tooling	Open-source	Web-agent benchmark execution and analysis	Best treated as an environment layer, not a full lifecycle eval platform	59
AgentDojo	Dynamic benchmark for prompt-injection attacks and defenses	Open-source	Adversarial evaluation	Best paired with a general harness rather than used alone	45

A useful selection rule is to choose one **primary system of record** for traces and experiments, then integrate specialized environments and red-team suites around it. For example, a team might use BrowserGym or OSWorld as the execution environment, Phoenix or LangSmith for traces, Promptfoo or DeepEval for rubric scoring, and AgentDojo for adversarial testing. The operational anti-pattern is tool sprawl without shared identifiers or trace schemas. OpenTelemetry's GenAI semantic conventions are promising here because they explicitly standardize spans, metrics, and events for GenAI and MCP-related systems. 60

Case studies from recent papers and industry

Anthropic's discussion of the Bolt AI team is one of the clearest public industry examples of an evaluation program growing from product pressure rather than academic benchmarking. Anthropic reports that, in

roughly three months, the team built an eval system that runs the agent, grades outputs with static analysis, uses browser agents to test apps, and employs LLM judges for behaviors such as instruction following. The lesson is that an effective pipeline can mix deterministic analysis, embodied end-to-end testing, and model-based grading rather than choosing one method. [61]

OpenAI's January 2026 guide on testing agent skills systematizes a similar but more explicitly engineering-oriented pattern. It frames an eval as "prompt → captured run → small set of checks → score," and recommends defining success in terms of outcome, process, style, and efficiency before writing the skill itself. This is a strong template for product teams because it turns agent behavior into a regression-tested contract instead of an informal demo. [62]

OSWorld is a valuable academic case because it made execution-based desktop evaluation mainstream for computer-use agents, but OSWorld-Verified is perhaps even more instructive. The original paper introduced a scalable environment spanning Ubuntu, Windows, and macOS with 369 real tasks and custom execution-based evaluation scripts. The later Verified release then focused on fixing numerous infrastructure and task-quality issues to improve signal quality and parallelization. The analytical lesson is that environment maintenance is part of the pipeline, and benchmark "ops work" can materially change conclusions. [63]

The BrowserGym ecosystem shows how research evaluation matures when environment standardization and experiment tooling are developed together. The BrowserGym paper motivated a unified gym-like environment with well-defined observation and action spaces, while the broader ecosystem added AgentLab to support agent creation, testing, and analysis across multiple benchmarks. That combination is important because it reduces one of the main sources of irreproducibility in agent research: every lab using slightly different wrappers, affordances, and logging conventions. [37]

τ-bench is a strong case study for reliability and human–agent interaction. Its key contribution is not just a more realistic conversational setup with users, policies, and tools, but a faithful evaluation method based on the end-of-conversation database state and the pass^k metric for repeated-trial reliability. The paper's finding that strong function-calling agents were still below 50% on success and below 25% on pass^8 in retail underscores why one-shot success claims are poor proxies for deployability. [41]

Anthropic's 2026 notes on infrastructure noise and BrowseComp eval awareness are case studies in **evaluation validity** rather than model capability. In one case, benchmark scores shifted by several percentage points depending on runtime resources and enforcement; in the other, the model identified the benchmark and found leaked or decryptable answers. Together, these episodes show that agent evaluation programs need contamination checks, environment normalization, and benchmark-integrity testing built into the process. [64]

Best practices and a modular reusable pipeline template

The best practices most strongly supported by the current evidence are straightforward but demanding in implementation. Define success before implementation. Keep outcome metrics and process metrics separate. Treat traces as first-class artifacts. Put deterministic checks wherever possible around environment state, tool arguments, and policy gates. Use LLM judges selectively and calibrate them with human review. Re-run tasks multiple times and report reliability, not just best-case pass rates. Version every dataset, grader, and environment. Make benchmark validity visible by logging runtime budgets, resource

ceilings, seeds, and environment hashes. These recommendations follow directly from OpenAI's evaluation guidance, Anthropic's eval vocabulary, LangSmith's lifecycle docs, and recent literature on reliability, contamination, and infrastructure noise. [65]

Recommended modular interfaces

A reusable pipeline is easiest to maintain when it exposes explicit interfaces instead of hard-coding benchmark-specific logic everywhere. The structure below is a recommended synthesis of Anthropic's task/trial/grader/transcript/outcome model, Inspect's dataset–solver–scorer task abstraction, OpenAI's traces/graders/datasets/eval-runs flow, and OpenTelemetry-style trace semantics. [66]

Interface	Responsibility	Minimal contract
`TaskSpec`	Defines what is being tested	input, allowed tools, success criteria, policy constraints, budget, split, provenance
`EnvironmentAdapter`	Creates and resets the world	reset, snapshot, restore, execute, export_artifacts
`AgentAdapter`	Wraps the agent framework under test	run(task, env, seed, limits) -> trace + outcome
`TraceStore`	Persists observability data	spans/events/messages/tool calls/artifacts keyed by run ID
`Grader`	Scores one aspect of behavior	grade(run) -> metric records + evidence spans
`Aggregator`	Computes experiment summaries	aggregate(metric records, metadata) -> dashboards and gates
`ReviewQueue`	Handles selective human inspection	enqueue(run, rubric), record_decision
`Gate`	Turns metrics into deployment decisions	pass/fail/defer with thresholds and confidence rules

Recommended data schemas

The following schemas are not standards; they are recommended templates designed to interoperate well with the sources above.

```
{
  "TaskSpec": {
    "task_id": "string",
    "domain": "string",
    "split": "dev|test|prod-sampled|redteam",
    "input": {"type": "object"},
    "success_criteria": [{"metric_key": "string", "operator": ">=",
```

```
"threshold": 0.0}],
    "policy_constraints": ["string"],
    "allowed_tools": ["string"],
    "environment_ref": "string",
    "budget": {"max_steps": 0, "max_tokens": 0, "max_runtime_sec": 0},
    "provenance": {"source": "manual|trace|mined|synthetic", "version":
"string"}
  }
}
```

```
{
  "TrialRecord": {
    "run_id": "string",
    "task_id": "string",
    "agent_version": "string",
    "framework": "string",
    "model": "string",
    "seed": 0,
    "environment_snapshot": "string",
    "started_at": "timestamp",
    "ended_at": "timestamp",
    "outcome": {
      "status": "success|fail|unsafe|timeout|error|needs_review",
      "environment_state_hash": "string",
      "final_artifacts": ["uri"]
    },
    "resource_usage": {
      "input_tokens": 0,
      "output_tokens": 0,
      "tool_calls": 0,
      "wall_time_ms": 0,
      "cost_usd": 0.0
    },
    "trace_uri": "string"
  }
}
```

```
{
  "GradeRecord": {
    "run_id": "string",
    "grader_id": "string",
    "target_scope": "run|step|tool_call|handoff|population",
    "metric_key": "string",
    "score": 0.0,
    "value": "optional categorical value",
```

```
    "confidence": 0.0,
    "threshold": "optional",
    "passed": true,
    "evidence": {
      "span_ids": ["string"],
      "artifact_uris": ["uri"]
    },
    "reviewer": "model|human|hybrid",
    "comment": "string"
  }
}
```

This schema shape supports several properties that matter in practice: replayability, evidence-backed grading, trace-to-metric joins, and straightforward transition from offline experiments to online monitoring. It also makes macro-eval analysis easier because lower-level labels can be aggregated without losing the connection to raw evidence. [67]

Automation points

A reusable program should automate evaluation at different cadences rather than trying to solve everything in one pass.

Cadence	What to run automatically	What usually remains manual
On commit or PR	deterministic checks, schema validation, small regression set, high-confidence safety rules	none, unless a gate fails
Nightly	repeat-run reliability, broader offline datasets, pairwise experiments, synthetic edge-case expansion	review of changed rubrics or surprise failures
Pre-release	long-horizon end-to-end env runs, adversarial suites, cost/latency sweeps, calibration checks	spot-check of LLM judge outputs and top failure clusters
Post-release continuous	online trace scoring, anomaly detection, sampled human review, incident triage	escalation of ambiguous or high-risk traces
Quarterly	benchmark refresh, leakage audit, environment maintenance, threshold recalibration	benchmark governance and policy review

This multi-cadence approach follows the development-to-production lifecycle described by LangSmith and OpenAI, the CI/CD patterns documented by Braintrust and DeepEval, and the macro-eval pattern shown in OpenAI's cookbook. [68]

Example implementation timeline

Because your scale, domain, and budget are unspecified, the timeline below is intentionally general and reusable. It assumes a lean but serious rollout rather than a frontier-lab-scale program.

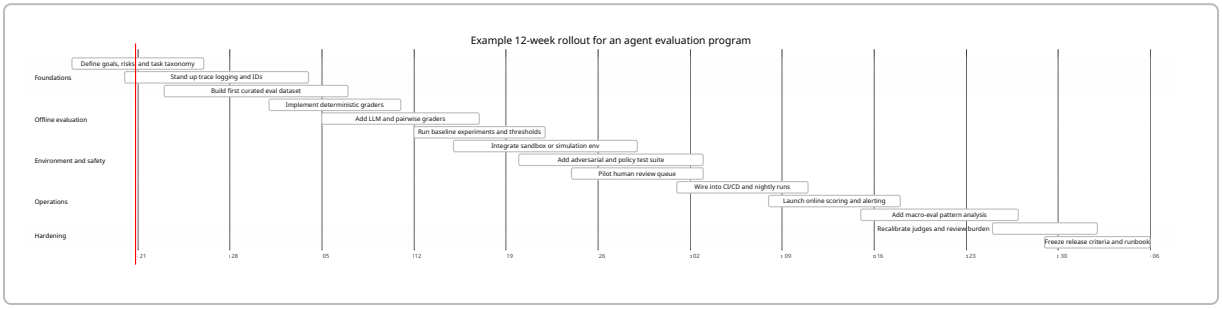

Deployment profiles when scale and budget are unspecified

If budget is tight, prefer a **lean** profile: curated offline dataset, deterministic checks, trace capture, and sampled human review. If volume is moderate, use a **standard** profile: add online trace scoring, reliability reruns, and environment-based end-to-end evals. If the domain is high-risk or the scale is large, use a **frontier** profile: add adversarial suites, macro-eval population analysis, judge calibration, contamination audits, and explicit environment governance. The progression is consistent with the staged lifecycle described by LangSmith, OpenAI, and NIST's risk-management framing. [69]

Limitations, open research questions, and ethical or regulatory considerations

The first major limitation is **benchmark validity**. Anthropic's BrowseComp write-up shows that web-enabled agents can encounter leaked answers or even infer that they are taking a benchmark, which threatens the interpretability of scores. This is not a niche concern; it means agent benchmarks increasingly need contamination audits, environment isolation, and replayable records of what the agent actually saw. [70]

The second limitation is **infrastructure confounding**. Anthropic's infrastructure-noise analysis shows that compute budgets and resource-enforcement choices can meaningfully change benchmark outcomes, sometimes by more than the gaps separating leading models. Agent evaluation therefore needs environment metadata, resource normalization, and sometimes score reporting conditioned on budget profiles rather than a single headline number. [71]

The third limitation is **judge reliability**. LLM-as-judge remains extremely useful, but the research literature increasingly emphasizes calibration and selective automation rather than blind trust. Recent work on automated grading recommends using confidence estimates to decide when to trust a model judge and when to escalate to humans; LangSmith and Langfuse operationalize similar ideas through annotation queues, review workflows, and mixed evaluator strategies. [72]

The fourth limitation is **Goodharting and reward hacking**. Recent work such as the Reward Hacking Benchmark, SpecBench, and EvilGenie shows that long-horizon agents can optimize visible tests or shallow

reward signals while missing the user's actual objective. For evaluation design, this means hidden checks, held-out validations, environment tamper detection, and orthogonal graders are becoming necessary even for seemingly "objective" coding tasks. [31]

Open research questions remain substantial. There is still no universal standard for measuring faithful versus merely plausible reasoning traces in agents. Interpretability at the trajectory level is improving, but current approaches remain fragmented across causal-faithfulness audits, trajectory-aware benchmarks, and internal-confidence methods. Similarly, multi-agent evaluation still lacks a stable consensus on how to score coordination quality, since milestone KPIs, collaboration graphs, and role-specialization metrics are young and highly benchmark-specific. [73]

The ethical and regulatory picture points toward more, not less, evaluation rigor. NIST says the AI RMF is intended to improve the incorporation of trustworthiness considerations into the design, development, use, and evaluation of AI systems, and its Generative AI Profile is meant to help organizations identify risks unique to generative AI. The EU AI Act requires continuous risk management, automatic record-keeping for high-risk AI systems, logging capabilities that support traceability and post-market monitoring, human oversight, and declared accuracy metrics for high-risk systems. These are directly relevant to agent-eval pipelines because agents create long traces, act through tools, and can produce high-impact downstream actions. [74]

Two additional governance sources reinforce the same direction. ISO/IEC 42001 frames AI management as a continuous management-system problem rather than a one-time model test, and the OECD AI Incidents and Hazards Monitor highlights the value of incident reporting and shared evidence about failure patterns. In practice, that means evaluation pipelines should be auditable, versioned, replayable, and connected to incident response and post-deployment monitoring rather than confined to model-launch week. [75]

The broad conclusion is that evaluation pipelines in agentic frameworks are becoming closer to **software reliability engineering plus safety engineering** than to classical NLP benchmarking. The most rigorous programs now combine benchmark realism, trace instrumentation, layered grading, controlled environments, reliability reruns, adversarial testing, calibrated human oversight, and continuous monitoring. That is the architecture most consistent with current primary papers, official platform guidance, and emerging regulatory expectations. [76]

---

[1] Agents • Cookbook
https://developers.openai.com/cookbook/topic/agents

[2] [4] [6] [12] [13] [17] [19] [23] [26] [33] [50] [65] [76] Evaluate agent workflows | OpenAI API
https://developers.openai.com/api/docs/guides/agent-evals

[3] [28] [64] [71] Quantifying infrastructure noise in agentic coding evals \ Anthropic
https://www.anthropic.com/engineering/infrastructure-noise

[5] [25] [66] Demystifying evals for AI agents \ Anthropic
https://www.anthropic.com/engineering/demystifying-evals-for-ai-agents

[7] [27] [62] Testing Agent Skills Systematically with Evals | OpenAI Developers
https://developers.openai.com/blog/eval-skills

8 9 14 Evaluation concepts - Docs by LangChain
https://docs.langchain.com/langsmith/evaluation-concepts

10 32 41 [2406.12045] $τ$-bench: A Benchmark for Tool-Agent-User Interaction in Real-World Domains
https://arxiv.org/abs/2406.12045

11 18 36 [2307.13854] WebArena: A Realistic Web Environment for Building Autonomous Agents
https://arxiv.org/abs/2307.13854

15 45 [2406.13352] AgentDojo: A Dynamic Environment to Evaluate Prompt ...
https://arxiv.org/abs/2406.13352?utm_source=chatgpt.com

16 How to run evaluations with pytest - Docs by LangChain
https://docs.langchain.com/langsmith/pytest

20 52 68 69 LangSmith Evaluation - Docs by LangChain
https://docs.langchain.com/langsmith/evaluation

21 Macro Evals for Agentic Systems
https://developers.openai.com/cookbook/examples/partners/macro_evals_for_agentic_systems/macro_evals_for_agentic_systems

22 24 44 [2510.04550] TRAJECT-Bench:A Trajectory-Aware Benchmark for Evaluating Agentic Tool Use
https://arxiv.org/abs/2510.04550

29 1.16. Probability calibration — scikit-learn 1.9.0 documentation
https://scikit-learn.org/stable/modules/calibration.html?utm_source=chatgpt.com

30 Part 1: Key Concepts in RL — Spinning Up documentation
https://spinningup.openai.com/en/latest/spinningup/rl_intro.html?utm_source=chatgpt.com

31 Reward Hacking Benchmark: Measuring Exploits in LLM Agents with Tool Use
https://arxiv.org/abs/2605.02964?utm_source=chatgpt.com

34 [2308.03688] AgentBench: Evaluating LLMs as Agents
https://arxiv.org/abs/2308.03688

35 48 [2311.12983] GAIA: a benchmark for General AI Assistants
https://arxiv.org/abs/2311.12983

37 [2412.05467] The BrowserGym Ecosystem for Web Agent Research
https://arxiv.org/abs/2412.05467

38 63 [2404.07972] OSWorld: Benchmarking Multimodal Agents for Open-Ended Tasks in Real Computer Environments
https://arxiv.org/abs/2404.07972

39 GitHub - xlang-ai/OSWorld: [NeurIPS 2024] OSWorld: Benchmarking Multimodal Agents for Open-Ended Tasks in Real Computer Environments · GitHub
https://github.com/xlang-ai/OSWorld

40 GitHub - SWE-bench/SWE-bench: SWE-bench: Can Language Models Resolve ...
https://github.com/swe-bench/SWE-bench?utm_source=chatgpt.com

42 [2412.14161] TheAgentCompany: Benchmarking LLM Agents on Consequential Real World Tasks
https://arxiv.org/abs/2412.14161

43 [2503.01935] MultiAgentBench: Evaluating the Collaboration and Competition of LLM agents
https://arxiv.org/abs/2503.01935

46 [2410.09024] AgentHarm: A Benchmark for Measuring Harmfulness of LLM Agents
https://arxiv.org/abs/2410.09024

47 [2506.14866] OS-Harm: A Benchmark for Measuring Safety of Computer Use Agents
https://arxiv.org/abs/2506.14866

49 51 Inspect
https://inspect.aisi.org.uk/

53 Evaluation of LLM Applications - Langfuse
https://langfuse.com/docs/evaluation/overview

54 GitHub - Arize-ai/phoenix: AI Observability & Evaluation · GitHub
https://github.com/Arize-ai/phoenix

55 Create experiments - Braintrust
https://www.braintrust.dev/docs/evaluate/run-evaluations

56 Introduction to DeepEval | DeepEval - The LLM Evaluation Framework
https://deepeval.com/docs/introduction

57 Evaluating Agents | MLflow AI Platform
https://mlflow.org/docs/latest/genai/eval-monitor/running-evaluation/agents/?utm_source=chatgpt.com

58 Intro | Promptfoo
https://www.promptfoo.dev/docs/intro/?utm_source=chatgpt.com

59 60 GitHub - ServiceNow/BrowserGym: BrowserGym, a Gym environment for web task automation · GitHub
https://github.com/ServiceNow/BrowserGym

61 Demystifying evals for AI agents \ Anthropic
https://www.anthropic.com/engineering/demystifying-evals-for-ai-agents?utm_source=chatgpt.com

67 Trace grading | OpenAI API
https://developers.openai.com/api/docs/guides/trace-grading

70 Eval awareness in Claude Opus 4.6’s BrowseComp performance \ Anthropic
https://www.anthropic.com/engineering/eval-awareness-browsecomp

72 When Can We Trust LLM Graders? Calibrating Confidence for Automated Assessment
https://arxiv.org/abs/2603.29559?utm_source=chatgpt.com

73 [2601.02314] Project Ariadne: A Structural Causal Framework for ...
https://arxiv.org/abs/2601.02314?utm_source=chatgpt.com

74 AI Risk Management Framework | NIST
https://www.nist.gov/itl/ai-risk-management-framework

75 ISO/IEC 42001:2023 - AI management systems
https://www.iso.org/standard/42001?utm_source=chatgpt.com